

Contents

Daily AI, BI & Research Briefing - 2026-05-10	1
Executive brief	1
What changed since yesterday	2
News radar	2
1. OpenAI: running Codex safely	2
2. OpenAI: GPT-5.5 trusted access for cyber	3
3. Microsoft Power BI: Execute DAX Queries REST API preview	3
4. Microsoft Power BI: translytical task flows	3
5. AWS: agent payments and workflow automation	4
Senior analyst insight	4
FMCG / sales & distribution practical implication	4
Academic paper brief: business value measurement	5
Main paper signal	5
Supporting academic signal	6
Worker experience and adoption mismatch	6
Technology learning bite: data quality / validation	6
Concept: semantic validation before AI action	6
AI engineering concept	7
Sandboxed agent execution	7
Daily reflection	7
Beginner German language drill	8
Vocabulary	8
Mini sentences	8
Practice prompt	8
Source quality review	8
Source weaknesses	10
Bottom line	10

Daily AI, BI & Research Briefing - 2026-05-10

Executive brief

- The strongest signal is not “better models”; it is the operationalisation layer around AI: sandboxing, approvals, telemetry, semantic models, task flows, and trusted execution.
- OpenAI’s Codex safety post and Microsoft Power BI’s DAX API / translytical task flow signals point in the same direction: AI and BI are moving closer to governed action, not just analysis.
- For business value measurement, the key question is whether AI shortens or improves real workflows without weakening control, auditability, data quality, or worker trust.
- The academic evidence remains useful but uneven: arXiv offers relevant early adoption signals, while OpenAlex/Crossref returned several noisy, future-dated, weakly matched, or duplicated items.
- Missing source categories remain a material weakness: official UK/regional

sources, FMCG/distribution sources, professional analyst sources, and reputable independent news are still underrepresented.

What changed since yesterday

Yesterday's memory entry framed the core bottleneck as organisational execution design: permissions, data quality, workflow fit, approvals, audit trails, and worker adoption.

Today's source pack reinforces that view. The strongest new continuity is that frontier AI and BI tooling are both becoming more operational:

- OpenAI's Codex safety post focuses on sandboxing, approvals, network policies, and agent telemetry. Source: <https://openai.com/index/running-codex-safely>
- Microsoft Power BI's Execute DAX Queries REST API preview points toward more programmable access to semantic models. Source: <https://community.fabric.microsoft.com/t5/BI-Updates-Blog/Execute-DAX-Queries-REST-API-Preview/ba-p/5177697>
- Power BI translytical task flows show reports becoming places where users can update records, add annotations, and trigger actions. Source: <https://community.fabric.microsoft.com/t5/BI-Updates-Blog/Translytical-Task-Flows-Generally-Available/ba-p/5173907>

Analyst interpretation: the adoption frontier is shifting from "Can the AI answer?" to "Can the organisation safely let the AI or BI workflow act?"

News radar

1. OpenAI: running Codex safely

Fact: OpenAI published a post on running Codex safely, highlighting sandboxing, approvals, network policies, and agent-native telemetry.

Source: <https://openai.com/index/running-codex-safely>

Why it matters: coding agents are not just productivity tools; they are execution systems. Once agents can modify code, call tools, or interact with repositories, the safety problem becomes operational governance.

Inference: organisations adopting coding agents will need controls similar to production change management: permission boundaries, audit logs, approval gates, and network restrictions.

Caution: this is a company primary source. It is valuable as a product and governance signal, but not neutral evidence of real-world effectiveness.

2. OpenAI: GPT-5.5 trusted access for cyber

Fact: OpenAI's source pack item says Trusted Access for Cyber is being expanded with GPT-5.5 and GPT-5.5-Cyber for verified defenders.

Source: <https://openai.com/index/gpt-5-5-with-trusted-access-for-cyber>

Why it matters: the cyber framing is important because it shows frontier AI access becoming segmented by user type, purpose, and verification status.

Inference: this pattern may spread to other high-risk enterprise domains: finance, healthcare, regulated operations, and critical infrastructure.

Caution: the briefing should not treat vendor claims about defender acceleration as independently validated unless external evaluation appears.

3. Microsoft Power BI: Execute DAX Queries REST API preview

Fact: Microsoft Power BI published an Execute DAX Queries REST API preview item on 2026-05-07.

Source: <https://community.fabric.microsoft.com/t5/Power-BI-Updates-Blog/Execute-DAX-Queries-REST-API-Preview/ba-p/5177697>

Why it matters: programmatic DAX execution increases the ability to integrate semantic model logic into automated workflows, apps, agents, and validation routines.

BI implication: the semantic model becomes less of a dashboard backend and more of a governed analytical service layer.

Practical risk: if semantic models are poorly governed, APIs can scale bad assumptions faster than manual reporting ever could.

4. Microsoft Power BI: translytical task flows

Fact: Power BI translytical task flows are described as enabling users to update records, add annotations, and trigger actions in external systems without leaving reports.

Source: <https://community.fabric.microsoft.com/t5/Power-BI-Updates-Blog/Translytical-Task-Flows-Generally-Available/ba-p/5173907>

Why it matters: this is a strong BI workflow-embedding signal. Reports are becoming operational interfaces.

Inference: for sales, distribution, service, and supply chain teams, the dashboard could become the control surface for exception handling: reviewing performance, correcting records, triggering follow-ups, and documenting decisions.

Control question: who is allowed to act from the report, under what conditions, with what validation, and with what audit trail?

5. AWS: agent payments and workflow automation

Fact: AWS's source pack includes Bedrock AgentCore Payments, developed with Coinbase and Stripe, and a Halliburton workflow-generation case using Amazon Bedrock.

Sources: - <https://aws.amazon.com/blogs/machine-learning/agents-that-transact-introducing-amazon-bedrock-agentcore-payments-built-with-coinbase-and-stripe/>
- <https://aws.amazon.com/blogs/machine-learning/halliburton-enhances-seismic-workflow-creation-with-amazon-bedrock-and-generative-ai/>

Why it matters: agentic systems are moving toward transaction and workflow execution, not just content generation.

Caution: AWS case studies are vendor-authored. They are useful for architecture patterns, but adoption claims and performance benefits should be treated as marketing-led unless independently validated.

Senior analyst insight

The market is converging on a new enterprise AI pattern:

AI creates or recommends actions; BI validates context; workflow systems execute; governance decides what is allowed.

That pattern is strategically more important than any single model announcement.

For organisations, the practical capability stack is becoming:

1. Clean data and semantic definitions.
2. Workflow-aware user interfaces.
3. Human approval gates where risk is material.
4. Auditability across AI suggestions and human decisions.
5. APIs that expose trusted metrics and business logic.
6. Monitoring for drift, misuse, and silent process failure.

The uncomfortable truth: many organisations are still trying to adopt AI on top of messy data, unclear ownership, weak process maps, and dashboards nobody acts on. In that environment, AI does not create transformation. It creates faster confusion.

FMCG / sales & distribution practical implication

For FMCG, sales, and distribution teams, the useful question is not "Can we add AI?" It is:

Which recurring decision loop can be shortened without losing commercial control?

Strong candidate workflows:

- Out-of-stock exception review.
- Distributor underperformance follow-up.
- Promotion compliance checks.
- Sales forecast variance explanation.
- Customer service escalation routing.
- Field sales visit prioritisation.
- SKU/customer profitability review.

A practical AI-BI workflow could look like this:

1. Power BI identifies a territory, customer, SKU, or route exception.
2. Semantic model logic confirms the metric definition.
3. AI explains the likely driver using approved data.
4. The user selects an action from a constrained menu.
5. A task flow updates a record, logs a note, or triggers a follow-up.
6. The action is auditable.

Senior judgement: this is where measurable value is more likely to appear: fewer handoffs, faster exception handling, cleaner follow-through, and better decision capture.

Risk: if master data, product hierarchies, customer mappings, or route data are weak, AI will amplify operational noise.

Academic paper brief: business value measurement

Main paper signal

Selected item: “When AI Meets Science: Research Diversity, Interdisciplinarity, Visibility, and Retractions across Disciplines in a Global Surge”

Source: <http://arxiv.org/abs/2605.06033v1>

Type: arXiv preprint / early signal, not settled evidence.

Why selected: Sunday’s academic rotation is business value measurement. This paper is not about SMEs directly, but it is relevant because it questions whether widespread AI adoption translates into deeper transformation or merely increased volume and visibility.

Reported source-pack claims:

- The paper studies AI adoption across countries and scientific domains.
- It reports substantial differences in timing and extent of adoption.
- It claims AI-supported works multiplied by at least four across all domains after 2015.
- It questions the transformative nature of the growth, arguing that AI-supported research remains concentrated in limited topics tied to computer science and conventional statistical frameworks.

- It also reports an association with citation premium and higher retraction rates.

Business interpretation: adoption volume is not the same as value creation. The SME equivalent would be counting AI tool usage, prompt volume, or licences as “success” while failing to measure whether decision quality, workflow speed, customer outcomes, or financial performance actually improved.

Research relevance: useful for sharpening the thesis distinction between: - adoption, - diffusion, - capability conversion, - measurable business value.

Caution: this is a preprint. It should be used as a conceptual warning signal, not as settled empirical proof.

Supporting academic signal

Worker experience and adoption mismatch

Selected item: “Making the Invisible Visible: Understanding the Mismatch Between Organizational Goals and Worker Experiences in AI Adoption”

Source: <http://arxiv.org/abs/2605.03078v1>

Why it matters: the paper directly supports the socio-technical adoption theme. The source pack says it identifies poor usability, interoperability issues, misaligned expectations, limited control, and insufficient communication as adoption barriers.

Use in research framing: this supports the argument that AI value depends on workflow fit and worker integration, not just executive sponsorship or tool availability.

Caution: also an arXiv preprint. Treat as early qualitative/theoretical signal.

Technology learning bite: data quality / validation

Today’s rotation points to data quality and validation.

Concept: semantic validation before AI action

A BI workflow should not let AI act directly on raw or ambiguous metrics. A safer pattern is:

1. Define the business metric in the semantic model.
2. Validate source freshness and completeness.
3. Check whether the metric has known caveats.
4. Let AI explain or recommend only after the metric passes validation.
5. Log the recommendation, user decision, and final action.

Example:

- Bad pattern: “AI, explain why sales dropped.”

- Better pattern: “Using the governed semantic model, explain sales variance for customers where data freshness is confirmed, returns are included, and customer hierarchy mapping is valid.”

This turns AI from a loose narrator into a controlled analytical assistant.

AI engineering concept

Sandboxed agent execution

A sandbox is a restricted environment where an AI agent can perform work without unrestricted access to the broader system.

In coding-agent or workflow-agent settings, the sandbox should control:

- Files the agent can read or modify.
- Network access.
- Tool permissions.
- Secrets and credentials.
- Approval requirements.
- Logs and telemetry.
- Rollback path.

Why it matters: as agents move from answering to acting, the risk moves from “wrong output” to “wrong action.”

Practical BI analogy: translytical BI workflows need a similar control mindset. If a dashboard can update records or trigger downstream actions, it needs permissions, validation, and audit logs.

Daily reflection

The strategic trap is mistaking adoption signals for value signals.

A company can have: - many AI users, - many dashboards, - many pilots, - many API integrations, - many automation ideas,

and still have weak business value if the workflow does not change or the decisions do not improve.

Today’s question:

Where is AI currently producing activity rather than measurable improvement?

A useful answer should name one workflow, one metric, one decision owner, and one evidence source.

Beginner German language drill

Theme: business value and workflow

Vocabulary

- der Wert = value
- der Prozess = process
- die Entscheidung = decision
- die Datenqualität = data quality
- messen = to measure
- verbessern = to improve
- prüfen = to check
- die Auswirkung = impact

Mini sentences

- Wir messen den Wert.
 - We measure the value.
- Die Datenqualität ist wichtig.
 - Data quality is important.
- Der Prozess muss klar sein.
 - The process must be clear.
- KI kann Entscheidungen unterstützen.
 - AI can support decisions.
- Wir prüfen die Auswirkung.
 - We check the impact.

Practice prompt

Translate into German:

1. We improve the process.
2. Data quality supports better decisions.
3. We measure the impact of AI.

Suggested answers:

1. Wir verbessern den Prozess.
 2. Datenqualität unterstützt bessere Entscheidungen.
 3. Wir messen die Auswirkung von KI.
-

Source quality review

Source group	Items used	Quality judgement	Analyst treatment
OpenAI News	Codex safety, trusted cyber access	Primary but not neutral	Strong product/governance signal; do not treat performance claims as independent evidence
Microsoft Power BI Blog	DAX API, translytical task flows	Vendor source but highly relevant for BI workflow direction	Strong practical tooling signal
AWS ML Blog	Agent payments, Halliburton workflow case	Vendor-authored, case-study style	Useful architecture signal; sceptical on claimed benefits
Google DeepMind Blog	AlphaEvolve and related AI transformation items	Primary frontier AI source	Relevant but less directly tied to BI/FMCG today
arXiv	AI adoption and business value papers	Preprint / early signal	Useful for research framing, not settled evidence
OpenAlex	SME/digital/AI metadata	Noisy	Several weak matches, future-dated items, and irrelevant concepts
Crossref	SME AI adoption/dynamic capabilities metadata	Mixed	Some relevant leads, but metadata is thin
Official UK/regional sources	Missing	Material weakness	Needed for SME, West Midlands, labour market, policy credibility
FMCG / retail / distribution sources	Missing	Material weakness	Needed for practical sector grounding
Professional analyst sources	Missing	Material weakness	Needed for independent strategy framing

Source weaknesses

The source pack is still skewed toward company blogs and academic metadata. This is acceptable for technical signal detection, but weak for senior strategy judgement.

Specific weaknesses:

- Microsoft AI Blog feed appears stale, returning 2022 items.
 - OpenAlex returned future-dated and weakly matched records.
 - Some OpenAlex items are irrelevant to SME AI adoption despite matching query terms.
 - Crossref items are potentially useful but too thin without abstracts or full context.
 - No official UK, ONS, DSIT, OECD, West Midlands, FMCG, retail, or logistics source appeared in the pack.
 - No independent news source was included to corroborate company announcements.
-

Bottom line

The briefing's strongest message is simple:

AI value will be won or lost in the governed workflow layer.

The next useful briefing should move from tooling signals to measurement discipline: how to prove that AI improves business outcomes rather than merely increasing activity.