

Contents

Daily AI, BI & Research Briefing — 2026-05-11	1
1. Executive brief	1
2. What changed since yesterday	2
3. News radar	2
3.1 Power BI: from reporting to operational interface	2
3.2 AWS and Halliburton: natural language to executable workflows	3
3.3 FMCG and supply chain: automation is entering scale-up mode	3
3.4 Trade and logistics: external shocks keep raising the value of faster decision loops	3
4. Senior analyst insight	4
5. FMCG / sales and distribution practical implication	4
6. Academic paper brief - SME AI adoption	5
7. Technology learning bite - SQL	5
8. AI engineering concept - function-based action layer	6
9. Daily reflection	6
10. Beginner German language drill	6
11. Source quality review	7
12. Missing source categories and next improvements	8
Selected source links	8

Daily AI, BI & Research Briefing — 2026-05-11

Generated from latest available source pack: data/source_packs/source_pack_2026-05-10.md

Scope: AI adoption, BI workflow embedding, FMCG / sales and distribution implications, and research-method continuity. This edition uses the latest available source pack rather than inventing a missing 2026-05-11 pack.

1. Executive brief

- **The strongest signal is decision-calibrated AI:** AI should not be used uniformly across every business question. Routine workflow decisions, operational exception handling, and high-stakes strategic choices need different levels of automation, human review, and evidence quality.
- **Power BI is moving deeper into the action layer:** translytical task flows let users update records, add annotations, and trigger external actions from reports; Dataflow Gen1 is now described as legacy while Gen2 is positioned as the future path for Fabric/Premium users. Source: Microsoft Power BI Blog.
- **FMCG and supply-chain evidence is more useful today than yesterday:** P&G, Iceland/invent.ai, Corvera, RedCloud, Amazon logistics, trade/tariff items, and UK logistics property signals all point to operational pressure around automation, replenishment, distribution capacity, and decision speed.
- **Vendor evidence is still not neutral evidence:** AWS reports a Halliburton proof-of-concept converting natural language into executable seismic workflows

with a claimed acceleration up to 95%, but the source is vendor-authored and should be treated as an architecture signal rather than independent proof.

- **The research gap remains live:** for Monday SME AI adoption, the source pack does not provide a strong peer-reviewed SME adoption paper. The best academic item is still an arXiv preprint on AI adoption in science, useful as a warning that adoption volume is not the same as transformation.

2. What changed since yesterday

Yesterday framed the bottleneck as governed execution: permissions, semantic validation, task flows, auditability, and human approval. Today adds a sharper layer: calibration by decision type. The question is no longer just whether AI can be embedded into a workflow; it is whether the organisation has matched the AI intervention to the risk, reversibility, and evidence requirements of that decision.

The MIT Sloan item on calibrating AI use to the decision at hand is the most useful professional framing. Its consumer-goods example separates decisions such as store-opening choices from broader brand-positioning choices. That distinction matters: some decisions are data-rich, repetitive, and testable; others are ambiguous, strategic, and politically loaded. Source: MIT Sloan Management Review.

The BI signal also matured. Translytical Power BI reports and Fabric User Data Functions imply dashboards are becoming operational interfaces. The Dataflow Gen1-to-Gen2 message adds a platform-modernisation warning: automation built on ageing data plumbing creates fragility, not transformation.

3. News radar

3.1 Power BI: from reporting to operational interface

Fact: The source pack includes Power BI April 2026 updates covering Copilot and AI, reporting, visuals, modelling, mobile experiences, performance, authoring workflows, and deprecation notices. It also includes translytical task flows as generally available and a Dataflows Gen1/Gen2 transition note. Sources: Power BI April update, Translytical task flows, Dataflows Gen1/Gen2.

Why it matters: BI is no longer just a read-only layer. When users can update records, add notes, and trigger downstream functions from a report, the report becomes a controlled business application.

Inference: for sales and distribution, the attractive use case is not another dashboard. It is an exception cockpit where a manager can see a customer/SKU/route issue, inspect the evidence, trigger a follow-up, and leave an auditable note without switching systems.

Caution: this comes from Microsoft sources. Treat it as a strong product-direction signal, not independent proof of customer ROI.

3.2 AWS and Halliburton: natural language to executable workflows

Fact: AWS describes a Halliburton proof-of-concept using Amazon Bedrock to convert natural-language queries into executable seismic workflows and answer questions over tool documentation. The source pack reports a claimed workflow acceleration of up to 95%. Source: AWS Machine Learning Blog.

Why it matters: the pattern is transferable even if the domain is not. Complex technical workflows often fail because users must know exactly which tool, parameter, and sequence to use. A natural-language layer can reduce friction if the action layer is constrained and validated.

Caution: this is a vendor-authored case. The 95% figure should be treated as a claimed proof-of-concept result, not a generalisable benchmark. The real lesson is architecture: connect domain knowledge, tool documentation, workflow generation, and validation.

3.3 FMCG and supply chain: automation is entering scale-up mode

Fact: Supply Chain Dive reports that P&G is shifting Supply Chain 3.0 and other platforms into large-scale rollout, with automation expected in warehouses and manufacturing plants. Source: Supply Chain Dive - P&G.

Fact: the source pack also includes retail/FMCG AI signals: Iceland partnering with invent.ai for inventory and replenishment operations, Corvera raising a further GBP 3m for an FMCG AI platform, and RedCloud promoting specialist FMCG AI agents trained on a claimed USD 6.9bn trade dataset. These items come through news-search results and should be verified against primary sources before being used as hard evidence.

Why it matters: FMCG adoption is moving toward replenishment, inventory, commercial decision support, and warehouse/manufacturing automation - exactly the operational zone where BI, data quality, and workflow design decide whether AI produces measurable value.

3.4 Trade and logistics: external shocks keep raising the value of faster decision loops

Fact: Supply Chain Dive reports a July 4 deadline around an EU tariff deal, a trade-court ruling on a 10% global tariff, and ongoing questions around Amazon Supply Chain Services as a logistics heavyweight. Sources: EU tariff deadline, 10% tariff ruling, Amazon Supply Chain Services.

UK logistics context: Logistics Manager reports Clearbell and Deva acquiring Banbury and West Byfleet assets. This is not an AI story, but Banbury-area logistics activity matters for local sales/distribution context. Source: Logistics Manager - Banbury and West Byfleet.

Inference: uncertainty around tariffs, capacity allocation, and logistics infrastructure strengthens the case for BI systems that can quickly show exposure by supplier, SKU, route, customer, and margin.

4. Senior analyst insight

Core judgement: the next layer of AI maturity is not general automation. It is decision-calibrated workflow design.

A blunt way to classify AI use:

- **Routine, reversible decisions:** automate more aggressively if data quality is high. Examples: formatting standard notes, routing low-risk tickets, suggesting replenishment checks.
- **Operational exception decisions:** use AI for diagnosis and recommendation, but keep a human action gate. Examples: out-of-stock risks, delivery variance, forecast exceptions, customer underperformance.
- **Commercial judgement decisions:** use AI as an analyst and challenger, not an autopilot. Examples: promotional strategy, customer segmentation, pricing moves, field-sales prioritisation.
- **Strategic or identity-level decisions:** use AI to structure options and expose assumptions; do not delegate judgement. Examples: brand repositioning, market entry, operating-model redesign.

The practical mistake is to treat all four as the same problem. They are not. The first category can be partly automated; the last category should be augmented, challenged, and documented.

For AI Business Intelligence work, this points toward a useful role definition: the analyst translates messy business questions into decision classes, data requirements, controls, and measurable outcome tests.

5. FMCG / sales and distribution practical implication

Best practical project angle: build an AI-assisted exception workflow around one sales/distribution decision loop, not a broad AI strategy deck.

Example workflow: customer/SKU replenishment risk review.

- Power BI flags customers, SKUs, or routes with abnormal sales movement, late replenishment, low inventory, or margin pressure.
- The semantic model defines the metric logic: actual sales, forecast, returns, inventory, delivery status, promotional period, and customer hierarchy.
- AI generates a short explanation, but only from approved fields and with missing-data caveats.
- The user chooses from controlled actions: contact account manager, create follow-up task, annotate exception, escalate supply issue, or mark as explained.
- The task flow writes the note/action back to the system and logs the user, timestamp, reason code, and data snapshot.

Why this is credible for FMCG: the current source pack has signals across P&G automation, inventory/replenishment AI, Power BI action flows, trade disruption, and logistics capacity. These are not random themes; they converge on operational decision speed under uncertainty.

What to measure: time to identify exceptions, time to action, reduction in unresolved exceptions, forecast error for selected SKU/customer groups, stockout frequency, duplicated manual work, and quality of decision notes.

Hard truth: if product master data, customer hierarchies, promotional calendars, or stock records are weak, AI will mostly produce confident commentary on bad inputs. Data quality remains the bottleneck.

6. Academic paper brief - SME AI adoption

Rotation theme: Monday is SME AI adoption. The retrieved source pack does not contain a strong peer-reviewed SME AI adoption paper suitable for today. Forcing a weak match would damage the briefing.

Selected academic signal: “When AI Meets Science: Research Diversity, Interdisciplinarity, Visibility, and Retractions across Disciplines in a Global Surge”. Source: arXiv 2605.06033v1.

Why use it: it is not an SME paper, but it is useful for the adoption-versus-transformation distinction. The source pack summary says the paper documents uneven AI adoption across countries and scientific domains, growth in AI-supported works after 2015, concentration around limited topics tied to computer science and statistical frameworks, and associations with citation premium and higher retraction rates.

Research interpretation: adoption volume is not automatically capability conversion. For SMEs, counting AI-tool usage, licences, prompts, or pilots is a weak success measure. Better measures ask whether AI changes decision speed, process quality, customer outcomes, resilience, or profitability.

Caution: this is an arXiv preprint and should be treated as an early signal, not settled evidence. Also, it is a proxy source for today because the source pack did not surface a stronger SME-specific academic item.

Research gap to keep: How do UK/West Midlands SMEs convert generative AI adoption into measurable business value once the first experimentation phase is over?

7. Technology learning bite - SQL

For AI/BI workflow work, SQL is not just a data extraction skill. It is how you define which records deserve attention before AI writes any explanation.

Mini concept: exception flags with CASE expressions.

```
SELECT
    customer_id,
    sku_id,
    sales_actual,
    sales_forecast,
    CASE
        WHEN sales_forecast = 0 THEN 'check_forecast_missing'
```

```

    WHEN sales_actual < sales_forecast * 0.75 THEN 'underperforming'
    WHEN sales_actual > sales_forecast * 1.25 THEN 'overperforming'
    ELSE 'normal'
END AS exception_flag
FROM customer_sku_sales;

```

Why it matters: AI should explain records after the exception logic is defined, not invent the exception logic on the fly. The analyst controls the rule; AI helps interpret and communicate it.

8. AI engineering concept - function-based action layer

A function-based action layer is a controlled set of actions that software or an AI assistant is allowed to call. In the Power BI example, reports connect to Fabric User Data Functions; in agentic systems, the same idea appears as tool/function calling.

The control principle is simple: do not give the AI unlimited system access. Give it narrow functions with defined inputs, validation, permissions, and logs.

- **Bad pattern:** “AI, fix the customer account issue” with broad system access.
- **Better pattern:** “AI, recommend one of five reason codes; user approves; approved function writes note and creates a follow-up task.”
- **Minimum safeguards:** input validation, role permissions, human approval for material actions, logging, rollback path, and monitoring for repeated error patterns.

This is where BI and AI engineering meet: semantic models define what is true; functions define what is allowed; governance defines who can act.

9. Daily reflection

Question: Which business decisions should AI be allowed to automate, which should it only support, and which should remain human-led?

A useful answer should name one workflow, one decision owner, one risk level, one data-quality requirement, and one measurable outcome.

A weak answer says “use AI to improve efficiency”. A strong answer says: “For customer/SKU replenishment exceptions, use AI to draft explanations and recommended next actions after SQL/semantic-model validation, with human approval before system updates, measured by time-to-action and unresolved exception rate.”

10. Beginner German language drill

Odd-date language rotation: German beginner.

Vocabulary

- die Entscheidung = decision
- der Arbeitsablauf = workflow
- die Lieferkette = supply chain

- die Datenqualitaet = data quality
- die Ausnahme = exception
- messen = to measure
- verbessern = to improve
- pruefen = to check

Mini sentences

- Wir pruefen die Datenqualitaet. - We check the data quality.
- Die Entscheidung ist wichtig. - The decision is important.
- KI verbessert den Arbeitsablauf nicht automatisch. - AI does not automatically improve the workflow.
- Wir messen die Auswirkung. - We measure the impact.

Practice

Translate into German:

1. We improve the workflow.
2. Data quality is important.
3. We measure the decision.

Suggested answers:

1. Wir verbessern den Arbeitsablauf.
2. Die Datenqualitaet ist wichtig.
3. Wir messen die Entscheidung.

11. Source quality review

Source group	Items used	Quality judgement	Treatment in briefing
Microsoft Power BI Blog	April update, translytical task flows, Dataflows Gen1/Gen2	Primary vendor source; strong product signal, not neutral ROI evidence.	Used for BI direction and workflow implications; claims handled cautiously.
AWS Machine Learning Blog	Halliburton Bedrock workflow proof-of-concept	Vendor-authored case; useful architecture signal; performance claim not generalisable.	Used to illustrate natural-language-to-workflow pattern with explicit caveat.
Supply Chain Dive / Logistics Manager	P&G, trade/tariff items, Amazon logistics, Banbury/West Byfleet assets	Specialist industry sources; useful context, but should be corroborated for legal/policy claims.	Used as operational pressure and FMCG/logistics context.

Source group	Items used	Quality judgement	Treatment in briefing
Google News search results	Iceland/invent.ai, Corvera, RedCloud, YouGov, Startups.co.uk	Thin summaries and RSS links; useful leads, weak as standalone evidence.	Used as signals only; not used for hard conclusions unless verified later.
MIT Sloan Management Review	Calibrate AI Use to the Decision at Hand	Professional insight; strong framing but not neutral academic evidence.	Used as the main decision-calibration lens.
arXiv	When AI Meets Science	Preprint; not peer reviewed; not SME-specific.	Used as early signal for adoption-vs-transformation caution.
Official UK / regional sources	GOV.UK query returned HTTP 422 error	Missing official evidence is a material weakness.	Flagged as source gap; not substituted with speculation.

12. Missing source categories and next improvements

- **Official UK/regional sources:** still weak today because the GOV.UK query errored. Add ONS, DSIT, West Midlands Combined Authority, OECD, Bank of England, EU AI Office, and Companies House where relevant.
- **Primary company sources for FMCG AI items:** verify Iceland/invent.ai, Corvera, and RedCloud from original company releases or credible trade sources before citing numbers heavily.
- **Peer-reviewed SME AI adoption literature:** Monday needs a stronger academic feed than arXiv/OpenAlex/Crossref snippets. Use Semantic Scholar, SSRN, CORE, and manual Google Scholar checks.
- **Independent analyst validation:** add Gartner/IDC/Forrester only if accessible, and balance consulting content with scepticism.

Bottom line: the strongest theme is not “AI everywhere”. It is AI matched to the decision, embedded in a governed workflow, and measured against operational outcomes.

Selected source links

- Power BI April 2026 Feature Summary
- Power BI Translytical Task Flows
- Power BI Dataflows Gen1 and Gen2
- AWS Halliburton Bedrock workflow case
- Supply Chain Dive - P&G Supply Chain 3.0
- Supply Chain Dive - EU tariff deadline
- Supply Chain Dive - 10% tariff ruling

- Supply Chain Dive - Amazon Supply Chain Services
- Logistics Manager - Banbury and West Byfleet assets
- MIT Sloan - Calibrate AI Use to the Decision at Hand
- arXiv - When AI Meets Science